


Interpretable Machine Learning

Shapley Values



Learning goals

- Learn what game theory is
- Understand the concept behind cooperative games
- Understand the Shapley value in game theory

COOPERATIVE GAMES IN GAME THEORY

► Shapley (1951)

- Game theory is the study of strategic games between players, “game” refers to any series of interactions between actors / agents with gains and losses of quantifiable utility value

COOPERATIVE GAMES IN GAME THEORY

► Shapley (1951)

- Game theory is the study of strategic games between players, “game” refers to any series of interactions between actors / agents with gains and losses of quantifiable utility value
- Cooperative games: For all possible players $P = \{1, \dots, p\}$, each subset of players $S \subseteq P$ forms a coalition – each coalition S achieves a certain payout

COOPERATIVE GAMES IN GAME THEORY

► Shapley (1951)

- Game theory is the study of strategic games between players, “game” refers to any series of interactions between actors / agents with gains and losses of quantifiable utility value
- Cooperative games: For all possible players $P = \{1, \dots, p\}$, each subset of players $S \subseteq P$ forms a coalition – each coalition S achieves a certain payout
- A value function $v : 2^{|P|} \mapsto \mathbb{R}$ maps all $2^{|P|}$ possible coalitions to their payout (or gain)

COOPERATIVE GAMES IN GAME THEORY

► Shapley (1951)

- Game theory is the study of strategic games between players, “game” refers to any series of interactions between actors / agents with gains and losses of quantifiable utility value
- Cooperative games: For all possible players $P = \{1, \dots, p\}$, each subset of players $S \subseteq P$ forms a coalition – each coalition S achieves a certain payout
- A value function $v : 2^{|P|} \mapsto \mathbb{R}$ maps all $2^{|P|}$ possible coalitions to their payout (or gain)
- $v(S)$ is the payout of coalition $S \subseteq P$ (payout of empty coalition must be zero: $v(\emptyset) = 0$)

COOPERATIVE GAMES IN GAME THEORY

► Shapley (1951)

- Game theory is the study of strategic games between players, “game” refers to any series of interactions between actors / agents with gains and losses of quantifiable utility value
- Cooperative games: For all possible players $P = \{1, \dots, p\}$, each subset of players $S \subseteq P$ forms a coalition – each coalition S achieves a certain payout
- A value function $v : 2^{|P|} \mapsto \mathbb{R}$ maps all $2^{|P|}$ possible coalitions to their payout (or gain)
- $v(S)$ is the payout of coalition $S \subseteq P$ (payout of empty coalition must be zero: $v(\emptyset) = 0$)
- As some players contribute more than others, we want to fairly divide the total achievable payout $v(P)$ among the players according to a player's individual contribution


COOPERATIVE GAMES IN GAME THEORY

► Shapley (1951)

- Game theory is the study of strategic games between players, “game” refers to any series of interactions between actors / agents with gains and losses of quantifiable utility value
- Cooperative games: For all possible players $P = \{1, \dots, p\}$, each subset of players $S \subseteq P$ forms a coalition – each coalition S achieves a certain payout
- A value function $v : 2^{|P|} \mapsto \mathbb{R}$ maps all $2^{|P|}$ possible coalitions to their payout (or gain)
- $v(S)$ is the payout of coalition $S \subseteq P$ (payout of empty coalition must be zero: $v(\emptyset) = 0$)
- As some players contribute more than others, we want to fairly divide the total achievable payout $v(P)$ among the players according to a player's individual contribution
- We call the individual payout per player $\phi_j, j \in P$ (later: Shapley value)


COOPERATIVE GAMES WITHOUT INTERACTIONS

Players do not interact
(payouts ☆ add up in each coalition)





What is a
fair payout
distribution?

Players do not interact





COOPERATIVE GAMES WITHOUT INTERACTIONS



⇒ Fair Payouts are Trivial Without Interactions

COOPERATIVE GAMES WITH INTERACTIONS




⇒ Unclear how to fairly distribute payouts when players interact

COOPERATIVE GAMES WITH INTERACTIONS

Question: What is a fair payout for player “yellow”?

Idea: Compute marginal contribution of the player of interest across different coalitions




- Compute the total payout of each coalition
- Compute difference in payouts for each coalition with and without player “yellow” (= marginal contribution)
- Average marginal contributions using appropriate weights

COOPERATIVE GAMES WITH INTERACTIONS


Question: What is a fair payout for player “yellow”?

Idea: Compute marginal contribution of the player of interest across different coalitions



Note: Each marginal contribution is weighted w.r.t. number of possible orders of its coalition

\leadsto More players in $S \Rightarrow$ more orderings of S



- Compute the total payout of each coalition
- Compute difference in payouts for each coalition with and without player “yellow” (= marginal contribution)
- Average marginal contributions using appropriate weights

SHAPLEY VALUE - SET DEFINITION

This idea refers to the **Shapley value** which assigns a payout value to each player according to its marginal contribution in all possible coalitions.

- Let $v(S \cup \{j\}) - v(S)$ be the marginal contribution of player j to coalition S
 \rightsquigarrow measures how much a player j increases the value of a coalition S

SHAPLEY VALUE - SET DEFINITION

This idea refers to the **Shapley value** which assigns a payout value to each player according to its marginal contribution in all possible coalitions.

- Let $v(S \cup \{j\}) - v(S)$ be the marginal contribution of player j to coalition S
 \rightsquigarrow measures how much a player j increases the value of a coalition S
- Average marginal contributions for all possible coalitions $S \subseteq P \setminus \{j\}$
 \rightsquigarrow order of how players join the coalition matters \Rightarrow different weights depending on size of S

SHAPLEY VALUE - SET DEFINITION

This idea refers to the **Shapley value** which assigns a payout value to each player according to its marginal contribution in all possible coalitions.

- Let $v(S \cup \{j\}) - v(S)$ be the marginal contribution of player j to coalition S
 \rightsquigarrow measures how much a player j increases the value of a coalition S
- Average marginal contributions for all possible coalitions $S \subseteq P \setminus \{j\}$
 \rightsquigarrow order of how players join the coalition matters \Rightarrow different weights depending on size of S
- Shapley value via **set definition** (weighting via multinomial coefficient):

$$\phi_j = \sum_{S \subseteq P \setminus \{j\}} \frac{|S|!(|P| - |S| - 1)!}{|P|!} (v(S \cup \{j\}) - v(S))$$

SHAPLEY VALUE - ORDER DEFINITION

The Shapley value was introduced as summation over sets $S \subseteq P \setminus \{j\}$, but it can be equivalently defined as a summation of all orders of players:

$$\phi_j = \frac{1}{|P|!} \sum_{\tau \in \Pi} (v(S_j^\tau \cup \{j\}) - v(S_j^\tau))$$

- Π : All possible orders of players (we have $|P|!$ in total)

SHAPLEY VALUE - ORDER DEFINITION

The Shapley value was introduced as summation over sets $S \subseteq P \setminus \{j\}$, but it can be equivalently defined as a summation of all orders of players:

$$\phi_j = \frac{1}{|P|!} \sum_{\tau \in \Pi} (v(S_j^\tau \cup \{j\}) - v(S_j^\tau))$$

- Π : All possible orders of players (we have $|P|!$ in total)
- S_j^τ : Set of players before player j in order $\tau = (\tau^{(1)}, \dots, \tau^{(p)})$ where $\tau^{(i)}$ is i -th element
 - \Rightarrow Example: Players 1, 2, 3 $\Rightarrow \Pi = \{(1, 2, 3), (1, 3, 2), (2, 1, 3), (2, 3, 1), (3, 1, 2), (3, 2, 1)\}$
 - \rightsquigarrow For order $\tau = (2, 1, 3)$ and player of interest $j = 3 \Rightarrow S_j^\tau = \{2, 1\}$
 - \rightsquigarrow For order $\tau = (3, 1, 2)$ and player of interest $j = 1 \Rightarrow S_j^\tau = \{3\}$
 - \rightsquigarrow For order $\tau = (3, 1, 2)$ and player of interest $j = 3 \Rightarrow S_j^\tau = \emptyset$

SHAPLEY VALUE - ORDER DEFINITION


The Shapley value was introduced as summation over sets $S \subseteq P \setminus \{j\}$, but it can be equivalently defined as a summation of all orders of players:

$$\phi_j = \frac{1}{|P|!} \sum_{\tau \in \Pi} (v(S_j^\tau \cup \{j\}) - v(S_j^\tau))$$

- Π : All possible orders of players (we have $|P|!$ in total)
- S_j^τ : Set of players before player j in order $\tau = (\tau^{(1)}, \dots, \tau^{(p)})$ where $\tau^{(i)}$ is i -th element
 - \Rightarrow Example: Players 1, 2, 3 $\Rightarrow \Pi = \{(\textcolor{red}{1}, \textcolor{red}{2}, 3), (1, 3, 2), (\textcolor{red}{2}, \textcolor{red}{1}, 3), (2, 3, 1), (3, 1, 2), (3, 2, 1)\}$
 - \rightsquigarrow For order $\tau = (2, 1, 3)$ and player of interest $j = 3 \Rightarrow S_j^\tau = \{2, 1\}$
 - \rightsquigarrow For order $\tau = (3, 1, 2)$ and player of interest $j = 1 \Rightarrow S_j^\tau = \{3\}$
 - \rightsquigarrow For order $\tau = (3, 1, 2)$ and player of interest $j = 3 \Rightarrow S_j^\tau = \emptyset$
- Order definition: Marginal contribution of orders that yield set $S = \{\textcolor{red}{1}, \textcolor{red}{2}\}$ is summed twice
 - \rightsquigarrow In set definition, it has the weight $\frac{2!(3-2-1)!}{3!} = \frac{2 \cdot 0!}{6} = \frac{2}{6}$



WEIGHTS FOR MARGINAL CONTRIBUTION - ILLUSTRATION

$|P|! = 6$ orders



WEIGHTS FOR MARGINAL CONTRIBUTION - ILLUSTRATION


$|P|! = 6$ orders




via orders

(players before "yellow" player)

$|S| = 2$
weight = $1/6$



$|S| = 2$
weight = $1/6$



$|S| = 1$
weight = $1/6$



$|S| = 1$
weight = $1/6$



$|S| = 0$
weight = $1/6$




$|S| = 0$
weight = $1/6$



WEIGHTS FOR MARGINAL CONTRIBUTION - ILLUSTRATION


$|P|! = 6$ orders




via orders

(players before "yellow" player)

$|S| = 2$
weight = $1/6$




$|S| = 2$
weight = $1/6$



$|S| = 1$
weight = $1/6$



$|S| = 1$
weight = $1/6$



$|S| = 0$
weight = $1/6$




$|S| = 0$
weight = $1/6$



via sets


$|S| = 2$
weight = $2/6$



$|S| = 1$
weight = $1/6$



$|S| = 1$
weight = $1/6$




$|S| = 0$
weight = $2/6$



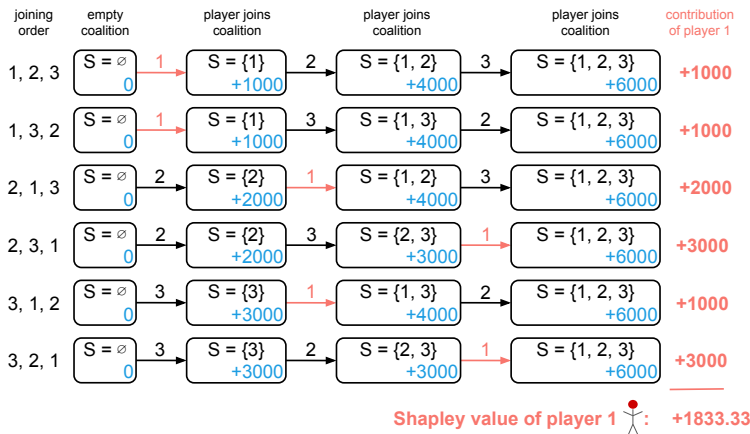
SHAPLEY VALUES - ILLUSTRATION

- Shapley value of player j is the marginal contribution to the value when it enters any coalition
- Produce all possible joining orders of player coalitions



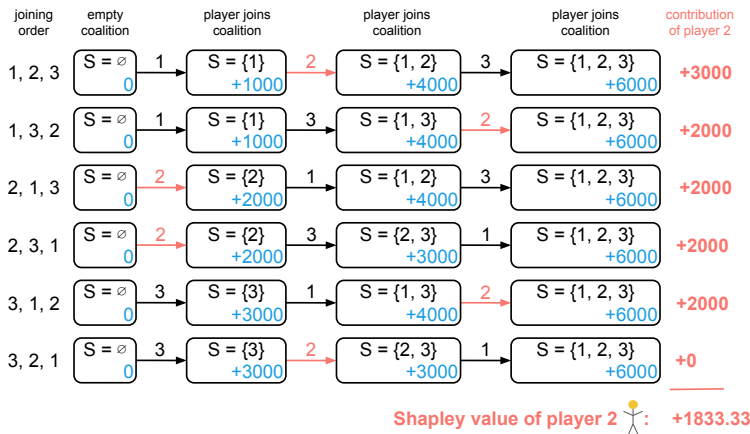
SHAPLEY VALUES - ILLUSTRATION

- Shapley value of player j is the marginal contribution to the value when it enters any coalition
- Produce all possible joining orders of player coalitions
- Measure and average the difference in payout after player 1 enters the coalition



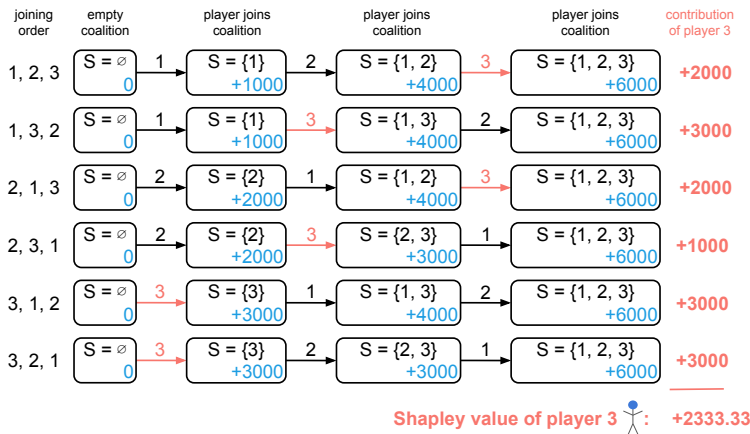
SHAPLEY VALUES - ILLUSTRATION

- Shapley value of player j is the marginal contribution to the value when it enters any coalition
- Produce all possible joining orders of player coalitions
- Measure and average the difference in payout after player 2 enters the coalition



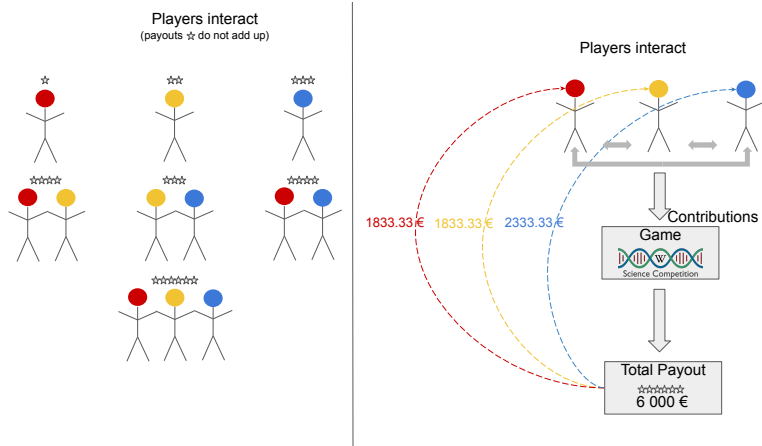
SHAPLEY VALUES - ILLUSTRATION

- Shapley value of player j is the marginal contribution to the value when it enters any coalition
- Produce all possible joining orders of player coalitions
- Measure and average the difference in payout after player 3 enters the coalition



SHAPLEY VALUES - ILLUSTRATION

- Shapley value of player j is the marginal contribution to the value when it enters any coalition
- Produce all possible joining orders of player coalitions



AXIOMS OF FAIR PAYOUTS

Why is this a fair payout solution?

One possibility to define fair payouts are the following axioms for a given value function v :

- **Efficiency:** Player contributions add up to the total payout of the game: $\sum_{j=1}^p \phi_j = v(P)$

AXIOMS OF FAIR PAYOUTS

Why is this a fair payout solution?

One possibility to define fair payouts are the following axioms for a given value function v :

- **Efficiency:** Player contributions add up to the total payout of the game: $\sum_{j=1}^p \phi_j = v(P)$
- **Symmetry:** Players $j, k \in P$ who contribute the same to any coalition get the same payout:
If $v(S \cup \{j\}) = v(S \cup \{k\})$ for all $S \subseteq P \setminus \{j, k\}$, then $\phi_j = \phi_k$

AXIOMS OF FAIR PAYOUTS

Why is this a fair payout solution?

One possibility to define fair payouts are the following axioms for a given value function v :

- **Efficiency:** Player contributions add up to the total payout of the game: $\sum_{j=1}^P \phi_j = v(P)$
- **Symmetry:** Players $j, k \in P$ who contribute the same to any coalition get the same payout:
If $v(S \cup \{j\}) = v(S \cup \{k\})$ for all $S \subseteq P \setminus \{j, k\}$, then $\phi_j = \phi_k$
- **Dummy/Null Player:** Payout is 0 for players who don't contribute to the value of any coalition:
If $v(S \cup \{j\}) = v(S) \quad \forall \quad S \subseteq P \setminus \{j\}$, then $\phi_j = 0$

AXIOMS OF FAIR PAYOUTS


Why is this a fair payout solution?

One possibility to define fair payouts are the following axioms for a given value function v :

- **Efficiency:** Player contributions add up to the total payout of the game: $\sum_{j=1}^P \phi_j = v(P)$
- **Symmetry:** Players $j, k \in P$ who contribute the same to any coalition get the same payout:
If $v(S \cup \{j\}) = v(S \cup \{k\})$ for all $S \subseteq P \setminus \{j, k\}$, then $\phi_j = \phi_k$
- **Dummy/Null Player:** Payout is 0 for players who don't contribute to the value of any coalition:
If $v(S \cup \{j\}) = v(S) \quad \forall \quad S \subseteq P \setminus \{j\}$, then $\phi_j = 0$
- **Additivity:** For a game v with combined payouts $v(S) = v_1(S) + v_2(S)$, the payout is the sum of payouts: $\phi_{j,v} = \phi_{j,v_1} + \phi_{j,v_2}$

Interpretable Machine Learning


Shapley Values for Local Explanations



Learning goals

- See model predictions as a cooperative game
- Transfer the Shapley value concept from game theory to machine learning

FROM GAME THEORY TO MACHINE LEARNING



FROM GAME THEORY TO MACHINE LEARNING

- Game: Make prediction $\hat{f}(x_1, x_2, \dots, x_p)$ for a single observation \mathbf{x}

FROM GAME THEORY TO MACHINE LEARNING


- Game: Make prediction $\hat{f}(x_1, x_2, \dots, x_p)$ for a single observation \mathbf{x}
- Players: Features $x_j, j \in \{1, \dots, p\}$ which cooperate to produce a prediction
 - ↪ How can we make a prediction with a subset of features without changing the model?
 - ↪ PD function: $\hat{f}_S(\mathbf{x}_S) := \int_{X_{-S}} \hat{f}(\mathbf{x}_S, X_{-S}) d\mathbb{P}_{X_{-S}}$ (“removing” by marginalizing over $-S$)

FROM GAME THEORY TO MACHINE LEARNING

- Game: Make prediction $\hat{f}(x_1, x_2, \dots, x_p)$ for a single observation \mathbf{x}
- Players: Features $x_j, j \in \{1, \dots, p\}$ which cooperate to produce a prediction
 - ↪ How can we make a prediction with a subset of features without changing the model?
 - ↪ PD function: $\hat{f}_S(\mathbf{x}_S) := \int_{\mathcal{X}_{-S}} \hat{f}(\mathbf{x}_S, \mathbf{x}_{-S}) d\mathbb{P}_{\mathcal{X}_{-S}}$ ("removing" by marginalizing over $-S$)
- Value function / payout of coalition $S \subseteq P$ for observation \mathbf{x} :

$$v(S) = \hat{f}_S(\mathbf{x}_S) - \mathbb{E}_{\mathbf{x}}(\hat{f}(\mathbf{x})), \text{ where } \hat{f}_S : \mathcal{X}_S \mapsto \mathcal{Y}$$

↪ subtraction of $\mathbb{E}_{\mathbf{x}}(\hat{f}(\mathbf{x}))$ ensures that v is a value function with $v(\emptyset) = 0$




FROM GAME THEORY TO MACHINE LEARNING

- Game: Make prediction $\hat{f}(x_1, x_2, \dots, x_p)$ for a single observation \mathbf{x}
- Players: Features $x_j, j \in \{1, \dots, p\}$ which cooperate to produce a prediction
 - \rightsquigarrow How can we make a prediction with a subset of features without changing the model?
 - \rightsquigarrow PD function: $\hat{f}_S(\mathbf{x}_S) := \int_{\mathcal{X}_{-S}} \hat{f}(\mathbf{x}_S, \mathbf{x}_{-S}) d\mathbb{P}_{\mathbf{x}_{-S}}$ ("removing" by marginalizing over $-S$)
- Value function / payout of coalition $S \subseteq P$ for observation \mathbf{x} :

$$v(S) = \hat{f}_S(\mathbf{x}_S) - \mathbb{E}_{\mathbf{x}}(\hat{f}(\mathbf{x})), \text{ where } \hat{f}_S : \mathcal{X}_S \mapsto \mathcal{Y}$$

\rightsquigarrow subtraction of $\mathbb{E}_{\mathbf{x}}(\hat{f}(\mathbf{x}))$ ensures that v is a value function with $v(\emptyset) = 0$



- Marginal contribution: $v(S \cup \{j\}) - v(S) = \hat{f}_{S \cup \{j\}}(\mathbf{x}_{S \cup \{j\}}) - \hat{f}_S(\mathbf{x}_S)$
 - $\rightsquigarrow \mathbb{E}_{\mathbf{x}}(\hat{f}(\mathbf{x}))$ cancels out due to the subtraction of value functions

SHAPLEY VALUE - DEFINITION

► Shapley (1953)

► Strumbelj et al. (2014)

Shapley value ϕ_j of feature j for observation \mathbf{x} via **order definition**:

$$\phi_j(\mathbf{x}) = \frac{1}{|P|!} \sum_{\tau \in \Pi} \underbrace{\hat{f}_{S_j^\tau \cup \{j\}}(\mathbf{x}_{S_j^\tau \cup \{j\}}) - \hat{f}_{S_j^\tau}(\mathbf{x}_{S_j^\tau})}_{\text{marginal contribution of feature } j}$$

- Interpretation: Feature x_j contributed ϕ_j to difference between $\hat{f}(\mathbf{x})$ and average prediction
 \rightsquigarrow Note: Marginal contributions and Shapley values can be negative
- For exact computation of $\phi_j(\mathbf{x})$, the PD function $\hat{f}_S(\mathbf{x}_S) = \frac{1}{n} \sum_{i=1}^n \hat{f}(\mathbf{x}_S, \mathbf{x}_{-S}^{(i)})$ for any set of features S can be used which yields

$$\phi_j(\mathbf{x}) = \frac{1}{|P|! \cdot n} \sum_{\tau \in \Pi} \sum_{i=1}^n \hat{f}(\mathbf{x}_{S_j^\tau \cup \{j\}}, \mathbf{x}_{-S_j^\tau \cup \{j\}}^{(i)}) - \hat{f}(\mathbf{x}_{S_j^\tau}, \mathbf{x}_{-S_j^\tau}^{(i)})$$

\rightsquigarrow Note: \hat{f}_S marginalizes over all other features $-S$ using all observations $i = 1, \dots, n$

ESTIMATION: A PRACTICAL PROBLEM

- Exact Shapley value computation is problematic for high-dimensional feature spaces
 - ↪ For 10 features, there are already $|P|! = 10! \approx 3.6$ million possible orders of features

ESTIMATION: A PRACTICAL PROBLEM

- Exact Shapley value computation is problematic for high-dimensional feature spaces
 \rightsquigarrow For 10 features, there are already $|P|! = 10! \approx 3.6$ million possible orders of features
- Additional problem due to estimation of the marginal prediction $\hat{f}_{S_j^\tau}$: Averaging over the entire data set for each coalition S_j^τ introduced by τ can be very expensive for large data sets

ESTIMATION: A PRACTICAL PROBLEM

- Exact Shapley value computation is problematic for high-dimensional feature spaces
 \rightsquigarrow For 10 features, there are already $|P|! = 10! \approx 3.6$ million possible orders of features
- Additional problem due to estimation of the marginal prediction $\hat{f}_{S_j^\tau}$: Averaging over the entire data set for each coalition S_j^τ introduced by τ can be very expensive for large data sets
- Solution to both problems is sampling: Instead of averaging over $|P|! \cdot n$ terms, we approximate it using a limited amount of M random samples of τ to build coalitions S_j^τ

ESTIMATION: A PRACTICAL PROBLEM

- Exact Shapley value computation is problematic for high-dimensional feature spaces
 - ↪ For 10 features, there are already $|P|! = 10! \approx 3.6$ million possible orders of features
- Additional problem due to estimation of the marginal prediction $\hat{f}_{S_j^\tau}$: Averaging over the entire data set for each coalition S_j^τ introduced by τ can be very expensive for large data sets
- Solution to both problems is sampling: Instead of averaging over $|P|! \cdot n$ terms, we approximate it using a limited amount of M random samples of τ to build coalitions S_j^τ
- M is a tradeoff between accuracy of the Shapley value and computational costs
 - ↪ The higher M , the closer to the exact Shapley values, but the more costly the computation

SHAPLEY VALUE APPROXIMATION - ILLUSTRATION

Definition

\mathbf{x} : obs. of interest

\mathbf{x} with feature values in S_m (other are replaced)

$$\phi_j(\mathbf{x}) = \frac{1}{M} \sum_{m=1}^M \left[\hat{f}(\mathbf{x}_{+j}^{(m)}) - \hat{f}(\mathbf{x}_{-j}^{(m)}) \right]$$

\mathbf{x} with feature values in $S_m \cup \{j\}$

	Temperature	Humidity	Windspeed	Year
\mathbf{x}	10.66	56	11	2012
\mathbf{x}_{+j}	10.66	56	random : $z_{windspeed}^{(m)}$	2012
\mathbf{x}_{-j}	10.66	56	random : $z_{windspeed}^{(m)}$	random : $z_{year}^{(m)}$

j

SHAPLEY VALUE APPROXIMATION - ILLUSTRATION

Definition

$$\phi_j(\mathbf{x}) = \frac{1}{M} \sum_{m=1}^M \underbrace{\left[\hat{f}(\mathbf{x}_{+j}^{(m)}) - \hat{f}(\mathbf{x}_{-j}^{(m)}) \right]}_{:= \Delta(j, S_m)}$$

Contribution of feature j
to coalition S_m

- $\Delta(j, S_m) = \hat{f}(\mathbf{x}_{+j}^{(m)}) - \hat{f}(\mathbf{x}_{-j}^{(m)})$ is the marginal contribution of feature j to coalition S_m
- Here: Feature *year* contributes +700 bike rentals if it joins coalition $S_m = \{\text{temp}, \text{hum}\}$

	Temperature	Humidity	Windspeed	Year	Count
\mathbf{x}	10.66	56	11	2012	
\mathbf{x}_{+j}	10.66	56	random : $z_{\text{windspeed}}^{(m)}$	2012	5600
\mathbf{x}_{-j}	10.66	56	random : $z_{\text{windspeed}}^{(m)}$	random : $z_{\text{year}}^{(m)}$	4900

j

\hat{f}

$\Delta(j, S_m)$
marginal contribution


SHAPLEY VALUE APPROXIMATION - ILLUSTRATION

Definition

average the contributions of feature j

$$\phi_j(\mathbf{x}) = \frac{1}{M} \sum_{m=1}^M \left[\hat{f}(\mathbf{x}_{+j}^{(m)}) - \hat{f}(\mathbf{x}_{-j}^{(m)}) \right]$$

- Compute marginal contribution of feature j towards the prediction across all randomly drawn feature coalitions S_1, \dots, S_m
- Average all M marginal contributions of feature j
- Shapley value ϕ_j is the payout of feature j , i.e., how much feature *year* contributed to the overall prediction in bicycle counts of a specific observation \mathbf{x}



REVISITED: AXIOMS FOR FAIR ATTRIBUTIONS

We take the general axioms for Shapley Values and apply it to predictions:

- **Efficiency:** Shapley values add up to the (centered) prediction: $\sum_{j=1}^p \phi_j = \hat{f}(\mathbf{x}) - \mathbb{E}_{\mathbf{x}}(\hat{f}(X))$

REVISITED: AXIOMS FOR FAIR ATTRIBUTIONS

We take the general axioms for Shapley Values and apply it to predictions:

- **Efficiency:** Shapley values add up to the (centered) prediction: $\sum_{j=1}^P \phi_j = \hat{f}(\mathbf{x}) - \mathbb{E}_{\mathbf{x}}(\hat{f}(X))$
- **Symmetry:** Two features j and k that contribute the same to the prediction get the same payout
 \rightsquigarrow interaction effects between features are fairly divided
 $\hat{f}_{S \cup \{j\}}(\mathbf{x}_{S \cup \{j\}}) = \hat{f}_{S \cup \{k\}}(\mathbf{x}_{S \cup \{k\}})$ for all $S \subseteq P \setminus \{j, k\}$ then $\phi_j = \phi_k$

REVISITED: AXIOMS FOR FAIR ATTRIBUTIONS

We take the general axioms for Shapley Values and apply it to predictions:


- **Efficiency:** Shapley values add up to the (centered) prediction: $\sum_{j=1}^P \phi_j = \hat{f}(\mathbf{x}) - \mathbb{E}_{\mathbf{x}}(\hat{f}(X))$
- **Symmetry:** Two features j and k that contribute the same to the prediction get the same payout
 \rightsquigarrow interaction effects between features are fairly divided
 $\hat{f}_{S \cup \{j\}}(\mathbf{x}_{S \cup \{j\}}) = \hat{f}_{S \cup \{k\}}(\mathbf{x}_{S \cup \{k\}})$ for all $S \subseteq P \setminus \{j, k\}$ then $\phi_j = \phi_k$
- **Dummy / Null Player:** Shapley value of a feature that does not influence the prediction is zero
 \rightsquigarrow if a feature was not selected by the model (e.g., tree or LASSO), its Shapley value is zero
 $\hat{f}_{S \cup \{j\}}(\mathbf{x}_{S \cup \{j\}}) = \hat{f}_S(\mathbf{x}_S)$ for all $S \subseteq P$ then $\phi_j = 0$

REVISITED: AXIOMS FOR FAIR ATTRIBUTIONS

We take the general axioms for Shapley Values and apply it to predictions:

- **Efficiency:** Shapley values add up to the (centered) prediction: $\sum_{j=1}^P \phi_j = \hat{f}(\mathbf{x}) - \mathbb{E}_{\mathbf{x}}(\hat{f}(X))$
- **Symmetry:** Two features j and k that contribute the same to the prediction get the same payout
 \rightsquigarrow interaction effects between features are fairly divided
 $\hat{f}_{S \cup \{j\}}(\mathbf{x}_{S \cup \{j\}}) = \hat{f}_{S \cup \{k\}}(\mathbf{x}_{S \cup \{k\}})$ for all $S \subseteq P \setminus \{j, k\}$ then $\phi_j = \phi_k$
- **Dummy / Null Player:** Shapley value of a feature that does not influence the prediction is zero
 \rightsquigarrow if a feature was not selected by the model (e.g., tree or LASSO), its Shapley value is zero
 $\hat{f}_{S \cup \{j\}}(\mathbf{x}_{S \cup \{j\}}) = \hat{f}_S(\mathbf{x}_S)$ for all $S \subseteq P$ then $\phi_j = 0$
- **Additivity:** For a prediction with combined payouts, the payout is the sum of payouts:
 $\phi_j(v_1) + \phi_j(v_2) \rightsquigarrow$ Shapley values for model ensembles can be combined

BIKE SHARING DATASET



- Shapley values of observation $i = 200$ from the bike sharing data
- Difference between model prediction of this observation and the average prediction of the data is fairly distributed among the features (i.e., $4434 - 4507 \approx -73$)
- Feature value temp = 28.5 has the most positive effect, with a contribution (increase of prediction) of about +400

ADVANTAGES AND DISADVANTAGES

Advantages:

- **Solid theoretical foundation** in game theory
- Prediction is **fairly distributed** among the feature values \rightsquigarrow easy to interpret for a user
- **Contrastive explanations** that compare the prediction with the average prediction

Disadvantages:

- Without sampling, Shapley values need a lot of computing time to inspect all possible coalitions
- Like many other IML methods, Shapley values suffer from the inclusion of unrealistic data observations when features are correlated

Interpretable Machine Learning

SHAP (SHapley Additive exPlanation) Values



Learning goals

- Get an intuition of additive feature attributions
- Understand the concept of Kernel SHAP
- Ability to interpret SHAP plots
- Global SHAP methods

FROM SHAPLEY TO SHAP


Remember: Shapley values explain the difference between actual and average prediction:

$$2573 - 4515 = 34 - 1654 - 323 = -1942$$

$$\hat{f}(\mathbf{x}) - \mathbb{E}(\hat{f}) = \phi_{hum} + \phi_{temp} + \phi_{ws}$$

↪ can be rewritten to

$$\hat{f}(\mathbf{x}) = \underbrace{\mathbb{E}(\hat{f})}_{\phi_0} + \phi_{hum} + \phi_{temp} + \phi_{ws}$$



SHAP DEFINITION

► Lundberg et al. 2017

Aim: Find an additive combination that explains the prediction of an observation \mathbf{x} by computing the contribution of each feature to the prediction using a (more efficient) estimation procedure.

Definition

- Define simplified (binary) coalition feature space $\mathbf{Z}' \in \{0, 1\}^{K \times p}$ with K rows and p columns
- Rows are referred to as $\mathbf{z}'^{(k)} = \{z_1'^{(k)}, \dots, z_p'^{(k)}\}$ with $k \in \{1, \dots, K\}$ (indexes k -th coalition)
- Columns are referred to as \mathbf{z}_j with $j \in \{1, \dots, p\}$ being the index of the original feature

Example:

Coalition	$\mathbf{z}'^{(k)}$	hum	temp	ws
\emptyset	$\mathbf{z}'^{(1)}$	0	0	0
hum	$\mathbf{z}'^{(2)}$	1	0	0
temp	$\mathbf{z}'^{(3)}$	0	1	0
ws	$\mathbf{z}'^{(4)}$	0	0	1
hum, temp	$\mathbf{z}'^{(5)}$	1	1	0
temp, ws	$\mathbf{z}'^{(6)}$	0	1	1
hum, ws	$\mathbf{z}'^{(7)}$	1	0	1
hum, temp, ws	$\mathbf{z}'^{(8)}$	1	1	1

SHAP DEFINITION


► Lundberg et al. 2017

Aim: Find an additive combination that explains the prediction of an observation \mathbf{x} by computing the contribution of each feature to the prediction using a (more efficient) estimation procedure.

Definition

- Define simplified (binary) coalition feature space $\mathbf{Z}' \in \{0, 1\}^{K \times p}$ with K rows and p columns
- Rows are referred to as $\mathbf{z}'^{(k)} = \{z_1'^{(k)}, \dots, z_p'^{(k)}\}$ with $k \in \{1, \dots, K\}$ (indexes k -th coalition)
- Columns are referred to as \mathbf{z}_j with $j \in \{1, \dots, p\}$ being the index of the original feature

$\mathbf{z}'^{(k)}$: Coalition
simplified features


$$g(\mathbf{z}'^{(k)}) = \phi_0 + \sum_{j=1}^p \phi_j z_j'^{(k)}$$

SHAP DEFINITION

► Lundberg et al. 2017

Aim: Find an additive combination that explains the prediction of an observation \mathbf{x} by computing the contribution of each feature to the prediction using a (more efficient) estimation procedure.

Definition

- Define simplified (binary) coalition feature space $\mathbf{Z}' \in \{0, 1\}^{K \times p}$ with K rows and p columns
- Rows are referred to as $\mathbf{z}'^{(k)} = \{z_1'^{(k)}, \dots, z_p'^{(k)}\}$ with $k \in \{1, \dots, K\}$ (indexes k -th coalition)
- Columns are referred to as \mathbf{z}_j with $j \in \{1, \dots, p\}$ being the index of the original feature

$\mathbf{z}'^{(k)}$: **Coalition**
simplified features

$$g(\mathbf{z}'^{(k)}) = \phi_0 + \sum_{j=1}^p \phi_j z_j'^{(k)}$$

ϕ_0 : **Null Output**
Average Model
Baseline ($\mathbb{E}(\hat{f})$)

SHAP DEFINITION

► Lundberg et al. 2017

Aim: Find an additive combination that explains the prediction of an observation \mathbf{x} by computing the contribution of each feature to the prediction using a (more efficient) estimation procedure.

Definition

- Define simplified (binary) coalition feature space $\mathbf{Z}' \in \{0, 1\}^{K \times p}$ with K rows and p columns
- Rows are referred to as $\mathbf{z}'^{(k)} = \{z_1'^{(k)}, \dots, z_p'^{(k)}\}$ with $k \in \{1, \dots, K\}$ (indexes k -th coalition)
- Columns are referred to as \mathbf{z}_j with $j \in \{1, \dots, p\}$ being the index of the original feature

$\mathbf{z}'^{(k)}$: **Coalition**
simplified features

$$g(\mathbf{z}'^{(k)}) = \phi_0 + \sum_{j=1}^p \phi_j z_j'^{(k)}$$

ϕ_0 : **Null Output**
Average Model
Baseline ($\mathbb{E}(\hat{f})$)

ϕ_j : **Attribution**
How much does
feature j change
the output for coal-
ition k

SHAP DEFINITION

► Lundberg et al. 2017

Aim: Find an additive combination that explains the prediction of an observation \mathbf{x} by computing the contribution of each feature to the prediction using a (more efficient) estimation procedure.

$g(\mathbf{z}'^{(k)})$: **Marginal Contribution**

Contribution of coalition $\mathbf{z}'^{(k)}$ to the prediction

$$g(\mathbf{z}'^{(k)}) = \phi_0 + \underbrace{\sum_{j=1}^p \phi_j z_j'^{(k)}}_{\text{Additive Feature Attribution}}$$

ϕ_j : **Shapley Values**

Additive Feature Attribution

Problem

How do we estimate the Shapley values ϕ_j ?

KERNEL SHAP - IN 5 STEPS

Definition: A kernel-based, model-agnostic method to compute Shapley values via local surrogate models (e.g. linear model)

- ➊ Sample coalitions
- ➋ Transfer coalitions into feature space & get predictions by applying ML model
- ➌ Compute weights through kernel
- ➍ Fit a weighted linear model
- ➎ Return Shapley values

KERNEL SHAP - IN 5 STEPS

Step 1: Sample coalitions

- Sample K coalitions from the simplified feature space

$$\mathbf{z}'^{(k)} \in \{0, 1\}^p, \quad k \in \{1, \dots, K\}$$


- For our simple example, we have in total $2^p = 2^3 = 8$ coalitions (without sampling)

Coalition	$\mathbf{z}'^{(k)}$	hum	temp	ws
\emptyset	$\mathbf{z}'^{(1)}$	0	0	0
hum	$\mathbf{z}'^{(2)}$	1	0	0
temp	$\mathbf{z}'^{(3)}$	0	1	0
ws	$\mathbf{z}'^{(4)}$	0	0	1
hum, temp	$\mathbf{z}'^{(5)}$	1	1	0
temp, ws	$\mathbf{z}'^{(6)}$	0	1	1
hum, ws	$\mathbf{z}'^{(7)}$	1	0	1
hum, temp, ws	$\mathbf{z}'^{(8)}$	1	1	1

KERNEL SHAP - IN 5 STEPS

Step 2: Transfer Coalitions into feature space & get predictions by applying ML model

- $\mathbf{z}'^{(k)}$ is 1 if features are part of the k -th coalition, 0 if they are absent
- To calculate predictions for these coalitions, we need to define a function which maps the binary feature space back to the original feature space



Coalition	$\mathbf{z}'^{(k)}$	hum	temp	ws	$\mathbf{x}^{coalition}$	hum	temp	ws
\emptyset	$\mathbf{z}'^{(1)}$	0	0	0	$\mathbf{x}^{\{\emptyset\}}$	\emptyset	\emptyset	\emptyset
hum	$\mathbf{z}'^{(2)}$	1	0	0	$\mathbf{x}^{\{hum\}}$	51.6	\emptyset	\emptyset
temp	$\mathbf{z}'^{(3)}$	0	1	0	$\mathbf{x}^{\{temp\}}$	\emptyset	5.1	\emptyset
ws	$\mathbf{z}'^{(4)}$	0	0	1	$\mathbf{x}^{\{ws\}}$	\emptyset	\emptyset	17.0
hum, temp	$\mathbf{z}'^{(5)}$	1	1	0	$\mathbf{x}^{\{hum,temp\}}$	51.6	5.1	\emptyset
temp, ws	$\mathbf{z}'^{(6)}$	0	1	1	$\mathbf{x}^{\{temp,ws\}}$	\emptyset	5.1	17.0
hum, ws	$\mathbf{z}'^{(7)}$	1	0	1	$\mathbf{x}^{\{hum,ws\}}$	51.6	\emptyset	17.0
hum, temp, ws	$\mathbf{z}'^{(8)}$	1	1	1	$\mathbf{x}^{\{hum,temp,ws\}}$	51.6	5.1	17.0

KERNEL SHAP - IN 5 STEPS

Step 2: Transfer Coalitions into feature space & get predictions by applying ML model

- Define $h_x(\mathbf{z}'^{(k)}) = \mathbf{z}^{(k)}$ where $h_x : \{0, 1\}^p \rightarrow \mathbb{R}^p$ maps 1's to feature values of observation \mathbf{x} for features part of the k -th coalition and 0's to feature values of a **randomly sampled observation** for features absent in the k -th coalition (feature values are permuted multiple times)
- Predict with ML model on this dataset $\hat{f} : \hat{f}(h_x(\mathbf{z}'^{(k)}))$

Coalition	$\mathbf{z}'^{(k)}$	$h_x(\mathbf{z}'^{(k)})$			$\mathbf{z}^{(k)}$	hum	temp	ws	$\hat{f}(h_x(\mathbf{z}'^{(k)}))$
\emptyset	$\mathbf{z}'^{(1)}$	0	0	0	$\mathbf{z}^{(1)}$	64.3	28.0	14.5	6211
hum	$\mathbf{z}'^{(2)}$	1	0	0	$\mathbf{z}^{(2)}$	51.6	28.0	14.5	5586
temp	$\mathbf{z}'^{(3)}$	0	1	0	$\mathbf{z}^{(3)}$	64.3	5.1	14.5	3295
ws	$\mathbf{z}'^{(4)}$	0	0	1	$\mathbf{z}^{(4)}$	64.3	28.0	17.0	5762
hum, temp	$\mathbf{z}'^{(5)}$	1	1	0	$\mathbf{z}^{(5)}$	51.6	5.1	14.5	2616
temp, ws	$\mathbf{z}'^{(6)}$	0	1	1	$\mathbf{z}^{(6)}$	64.3	5.1	17.0	2900
hum, ws	$\mathbf{z}'^{(7)}$	1	0	1	$\mathbf{z}^{(7)}$	51.6	28.0	17.0	5411
hum, temp, ws	$\mathbf{z}'^{(8)}$	1	1	1	$\mathbf{z}^{(8)}$	51.6	5.1	17.0	2573

KERNEL SHAP - IN 5 STEPS

Step 3: Compute weights through Kernel

Intuition: We learn most about individual features if we can study their effects in isolation or at maximal interaction: Small coalitions (few 1's) and large coalitions (i.e. many 1's) get the largest weights



KERNEL SHAP - IN 5 STEPS

Step 3: Compute weights through Kernel [▶ see shapley_kernel_proof.pdf](#)

Intuition: We learn most about individual features if we can study their effects in isolation or at maximal interaction: Small coalitions (few 1's) and large coalitions (i.e. many 1's) get the largest weights

The diagram illustrates the formula for the kernel weight $\pi_x(\mathbf{z}'^{(k)})$. The formula is:

$$\pi_x(\mathbf{z}'^{(k)}) = \frac{(p-1)}{\binom{p}{|\mathbf{z}'^{(k)}|} |\mathbf{z}'^{(k)}| (p - |\mathbf{z}'^{(k)}|)}$$

Annotations:

- $\pi_x(\mathbf{z}'^{(k)})$: kernel weight for coalition $\mathbf{z}'^{(k)}$
- p : Number of features in \mathbf{x}
- $|\mathbf{z}'^{(k)}|$: coalition size / sum of 1s in $\mathbf{z}'^{(k)}$

KERNEL SHAP - IN 5 STEPS

Step 3: Compute weights through Kernel

Purpose: to include this knowledge in the local surrogate model (linear regression), we calculate weights for each coalition which are the observations of the linear regression

$$\pi_x(\mathbf{z}') = \frac{(p-1)}{\binom{p}{|\mathbf{z}'|} |\mathbf{z}'| (p-|\mathbf{z}'|)} \rightsquigarrow \pi_x(\mathbf{z}' = (1, 0, 0)) = \frac{(3-1)}{\binom{3}{1} 1 (3-1)} = \frac{1}{3}$$

Coalition	$\mathbf{z}'^{(k)}$	hum	temp	ws	weight
\emptyset	$\mathbf{z}'^{(1)}$	0	0	0	∞
hum	$\mathbf{z}'^{(2)}$	1	0	0	0.33
temp	$\mathbf{z}'^{(3)}$	0	1	0	0.33
ws	$\mathbf{z}'^{(4)}$	0	0	1	0.33
hum, temp	$\mathbf{z}'^{(5)}$	1	1	0	0.33
temp, ws	$\mathbf{z}'^{(6)}$	0	1	1	0.33
hum, ws	$\mathbf{z}'^{(7)}$	1	0	1	0.33
hum, temp, ws	$\mathbf{z}'^{(8)}$	1	1	1	∞

KERNEL SHAP - IN 5 STEPS

Step 3: Compute weights through Kernel

Purpose: to include this knowledge in the local surrogate model (linear regression), we calculate weights for each coalition which are the observations of the linear regression

Coalition	$\mathbf{z}'^{(k)}$	hum	temp	ws	weight
\emptyset	$\mathbf{z}'^{(1)}$	0	0	0	∞
hum	$\mathbf{z}'^{(2)}$	1	0	0	0.33
temp	$\mathbf{z}'^{(3)}$	0	1	0	0.33
ws	$\mathbf{z}'^{(4)}$	0	0	1	0.33
hum, temp	$\mathbf{z}'^{(5)}$	1	1	0	0.33
temp, ws	$\mathbf{z}'^{(6)}$	0	1	1	0.33
hum, ws	$\mathbf{z}'^{(7)}$	1	0	1	0.33
hum, temp, ws	$\mathbf{z}'^{(8)}$	1	1	1	∞

↪ weights for empty and full set are infinity and not used as observations for the linear regression

↪ instead constraints are used such that properties (local accuracy and missingness) are satisfied

KERNEL SHAP - IN 5 STEPS

Step 4: Fit a weighted linear model

Aim: Estimate a weighted linear model with Shapley values being the coefficients ϕ_j

$$g\left(\mathbf{z}'^{(k)}\right)=\phi_0+\sum_{j=1}^p\phi_jz_j'^{(k)}$$

and minimize by WLS using the weights π_x of step 3

$$L\left(\hat{f},g,\pi_x\right)=\sum_{k=1}^K\left[\hat{f}\left(h_x\left(\mathbf{z}'^{(k)}\right)\right)-g\left(\mathbf{z}'^{(k)}\right)\right]^2\pi_x\left(\mathbf{z}'^{(k)}\right)$$

with $\phi_0=\mathbb{E}(\hat{f})$ and $\phi_p=\hat{f}(x)-\sum_{j=0}^{p-1}\phi_j$ we receive a $p-1$ dimensional linear regression problem


KERNEL SHAP - IN 5 STEPS


Step 4: Fit a weighted linear model

Aim: Estimate a weighted linear model with Shapley values being the coefficients ϕ_j

$$g(\mathbf{z}'^{(k)}) = \phi_0 + \sum_{j=1}^p \phi_j z_j'^{(k)} \rightsquigarrow g(\mathbf{z}'^{(k)}) = 4515 + 34 \cdot z_1'^{(k)} - 1654 \cdot z_2'^{(k)} - 323 \cdot z_3'^{(k)}$$

$\mathbf{z}'^{(k)}$	hum	temp	ws	weight	\hat{f}
$\mathbf{z}'^{(2)}$	1	0	0	0.33	4635
$\mathbf{z}'^{(3)}$	0	1	0	0.33	3087
$\mathbf{z}'^{(4)}$	0	0	1	0.33	4359
$\mathbf{z}'^{(5)}$	1	1	0	0.33	3060
$\mathbf{z}'^{(6)}$	0	1	1	0.33	2623
$\mathbf{z}'^{(7)}$	1	0	1	0.33	4450


input



output

KERNEL SHAP - IN 5 STEPS

Step 5: Return SHAP values

Intuition: Estimated Kernel SHAP values are equivalent to Shapley values

$$g(\mathbf{z}'^{(8)}) = \hat{f}(h_x(\mathbf{z}'^{(8)})) = 4515 + 34 \cdot 1 - 1654 \cdot 1 - 323 \cdot 1 = \underbrace{\mathbb{E}(\hat{f})}_{\phi_0} + \phi_{hum} + \phi_{temp} + \phi_{ws} = \hat{f}(\mathbf{x}) = 2573$$



PROPERTIES

Local Accuracy

$$f(\mathbf{x}) = g(\mathbf{x}') = \phi_0 + \sum_{j=1}^p \phi_j x'_j$$

Intuition: If the coalition includes all features ($\mathbf{x}' \in \{1\}^p$), the attributions ϕ_j and the null output ϕ_0 sum up to the original model output $f(\mathbf{x})$

Local accuracy corresponds to the **axiom of efficiency** in Shapley game theory

PROPERTIES

Local Accuracy

$$f(\mathbf{x}) = g(\mathbf{x}') = \phi_0 + \sum_{j=1}^p \phi_j x'_j$$

Missingness

$$x'_j = 0 \implies \phi_j = 0$$

Intuition: A missing feature gets an attribution of zero

PROPERTIES

Local Accuracy

$$f(\mathbf{x}) = g(\mathbf{x}') = \phi_0 + \sum_{j=1}^p \phi_j x'_j$$

Missingness

$$x'_j = 0 \implies \phi_j = 0$$

Consistency

$\hat{f}_x(\mathbf{z}'^{(k)}) = \hat{f}(h_x(\mathbf{z}'^{(k)}))$ and $\mathbf{z}'_{-j}{}^{(k)}$ denote setting $z_j'^{(k)} = 0$. For any two models \hat{f} and \hat{f}' , if

$$\hat{f}'_x(\mathbf{z}'^{(k)}) - \hat{f}'_x(\mathbf{z}'_{-j}{}^{(k)}) \geq \hat{f}_x(\mathbf{z}'^{(k)}) - \hat{f}_x(\mathbf{z}'_{-j}{}^{(k)})$$

for all inputs $\mathbf{z}'^{(k)} \in \{0, 1\}^p$, then

$$\phi_j(\hat{f}', \mathbf{x}) \geq \phi_j(\hat{f}, \mathbf{x})$$

PROPERTIES

Local Accuracy

$$f(\mathbf{x}) = g(\mathbf{x}') = \phi_0 + \sum_{j=1}^p \phi_j x'_j$$

Missingness

$$x'_j = 0 \implies \phi_j = 0$$

Consistency


$$\hat{f}'_x(\mathbf{z}'^{(k)}) - \hat{f}'_x(\mathbf{z}'^{(k)}_{-j}) \geq \hat{f}_x(\mathbf{z}'^{(k)}) - \hat{f}_x(\mathbf{z}'^{(k)}_{-j}) \implies \phi_j(\hat{f}', \mathbf{x}) \geq \phi_j(\hat{f}, \mathbf{x})$$

Intuition: If a model changes so that the marginal contribution of a feature value increases or stays the same, the Shapley value also increases or stays the same

From **consistency** the Shapley **axioms of additivity, dummy and symmetry** follow

Interpretable Machine Learning

Global SHAP



Learning goals

- Get an intuition of additive feature attributions
- Understand the concept of Kernel SHAP
- Ability to interpret SHAP plots
- Global SHAP methods

Idea:


- Run SHAP for every observation and thereby get a matrix of Shapley values
- The matrix has one row per data observation and one column per feature
- We can interpret the model globally by analyzing the Shapley values in this matrix

$$\Phi = \begin{bmatrix} \phi_{11} & \phi_{12} & \phi_{13} & \dots & \phi_{1p} \\ \phi_{21} & \phi_{22} & \phi_{23} & \dots & \phi_{2p} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ \phi_{n1} & \phi_{n2} & \phi_{n3} & \dots & \phi_{np} \end{bmatrix}$$

FEATURE IMPORTANCE

Idea: Average the absolute Shapley values of each feature over all observations. This corresponds to calculating averages column by column in Φ

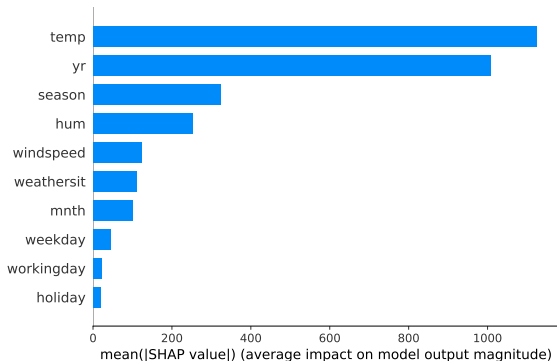
$$I_j = \frac{1}{n} \sum_{i=1}^n |\phi_j^{(i)}|$$



FEATURE IMPORTANCE

Interpretation:


- The features temperature and year have by far the highest influence on the model's prediction
- Compared to Shapley values, no effect direction is provided, but instead a feature ranking similar to PFI
- However, Shapley FI is based on the model's predictions only while PFI is based on the model's performance (loss)



SUMMARY PLOT

Combines feature importance with feature effects


- Each point is a Shapley value for a feature and an observation
- The color represents the value of the feature from low to high
- Overlapping points are jittered in y-axis direction



SUMMARY PLOT


Interpretation:

- Low temperatures have a negative impact while high temperatures lead to more bike rentals
- Year: two point clouds for 2011 and 2012 (other categorical features are gray)
- A high humidity has a huge, negative impact on the bike rental, while low humidity has a rather minor positive impact on bike rentals



DEPENDENCE PLOT


- Visualize the marginal contribution of a feature similar to the PDP
- Plot a point with the feature value on the x-axis and the corresponding Shapley value on the y-axis



DEPENDENCE PLOT

Interpretation:


- Increasing temperatures induce increasing bike rentals until 25°C
- If it gets too hot, the bike rentals decrease



DEPENDENCE PLOT

Interpretation:

- We can colour the observations by a second feature to detect interactions
- Visibly the temperatures interaction with the season is very strong



DISCUSSION

Advantages

- All the advantages of Shapley values
- Unify the field of interpretable machine learning in the class of additive feature attribution methods
- Has a fast implementation for tree-based models
- Various global interpretation methods

Disadvantages

- Disadvantages of Shapley values also apply to SHAP
- KernelSHAP is slow (TreeSHAP can be used as a faster alternative for tree-based models
▶ Lundberg et al 2018 – and for an intuitive explanation ▶ see Sukumar: TreeSHAP)
- KernelSHAP ignores feature dependence